



Developers

Welcome

Introduction

Getting Started

Models and Pricing

Rate Limits

■ Provisioned Throughput new

Regional Endpoints

Debugging Errors

Docs MCP

Release Notes

Model Capabilities

> Text

> Images

> Video new

> Audio new

> Files

> Legacy

Files & Collections

Files Overview

Managing Files

Collections

Collections via API

Collection Metadata

🖥️ API Console

● Loading...

Enterprise

Provisioned Throughput

Copy for LLM | View as Markdown

Provisioned Throughput allows enterprise customers to purchase dedicated input and output token capacity for specific models. Buy units with a minimum 30-day commitment for predictable, high-performance API access.



Contact support@x.ai to get started with Provisioned Throughput.

Key Benefits

- **Predictable latency** — Faster, more consistent response times compared to pay-as-you-go, even during peak usage
- **Uncapped scale** — Your purchased capacity adds directly to your rate limits; overages use standard pay-as-you-go rates
- **High reliability** — 99.9% uptime SLA with enterprise-grade availability guarantees

Pricing

Each unit costs **\$10.00 per day** and provides a fixed amount of tokens per minute (TPM):