



Developers

- Welcome
- Introduction
- Getting Started
- Models and Pricing
- Rate Limits
- Provisioned Throughput new
- Regional Endpoints
- Debugging Errors
- Docs MCP
- Release Notes

Model Capabilities

- > Text
- > Images
- > Video new
- > Audio new
- > Files
- > Legacy

Files & Collections

- Files Overview
- Managing Files
- Collections
- Collections via API
- Collection Metadata

API Console

Enterprise

Provisioned Throughput

Copy for LLM | View as Markdown

Provisioned Throughput allows enterprise customers to purchase dedicated input and output token capacity for specific models. Buy units with a minimum 30-day commitment for predictable, high-performance API access.



Contact support@x.ai to get started with Provisioned Throughput.

Key Benefits

- **Predictable latency** — Faster, more consistent response times compared to pay-as-you-go, even during peak usage
- **Uncapped scale** — Your purchased capacity adds directly to your rate limits; overages use standard pay-as-you-go rates
- **High reliability** — 99.9% uptime SLA with enterprise-grade availability guarantees

Pricing

Each unit costs **\$10.00 per day** and provides a fixed amount of tokens per minute (TPM):